

I S H A

Software Engineer · AI & Backend Systems

Noida, India · +91 98556 94722 · techbyisha@gmail.com · [LinkedIn](#) · [GitHub](#) · [Portfolio](#)

PROFESSIONAL SUMMARY

Backend & AI Engineer with nearly 2 years of production experience across fintech, insurance, and enterprise analytics. Specialises in RAG pipelines, knowledge graph architectures, and Node.js microservices that directly impact business KPIs. Built systems processing 50K+ record datasets, achieving sub-3s LLM query latency and 99.9% uptime on multi-model deployments. Seeking a high-ownership AI engineering role at a global MNC.

TECHNICAL SKILLS

AI / GenAI LangChain · RAG Pipelines · LLM Deployments · OpenAI API · FAISS · Pinecone · Neo4j (Cypher) · FastAPI · Model Context Protocol (MCP)

Languages Python · JavaScript (ES6+) · SQL · C++

Backend Node.js · Express.js · REST APIs · JWT Auth · Microservices

Data & Analytics Data Pipelines · Data Profiling · KPI Design · Excel Automation · Power BI · Streamlit

Cloud & Infra Google Cloud · Firebase · PostgreSQL

Automation n8n Workflow Automation · Web Scraping · Google Sheets Automation · API Integrations

Tools Git · Jira · Postman

EXPERIENCE

Data Analyst – Knowledge Initiative (Contract) · **Synopsys** · Noida, India *Feb 2026 – Apr 2026*

- Engineered Python data pipelines ingesting 50K+ enterprise knowledge records, cutting manual processing time by ~70% and enabling weekly automated reporting cycles.
- Designed a 3-KPI governance framework (Drop-off Rate, Knowledge Efficiency Ratio, YoY Growth) adopted by global engineering and knowledge teams to prioritise content investment.
- Built automated duplicate-detection scripts that flagged 8–12% of records as redundant, directly improving search precision and content discoverability scores.
- Delivered executive analytics dashboards surfacing actionable insights on content lifecycle gaps, informing Q1 2026 knowledge governance roadmap decisions.

Software Engineer II · **ChatFin AI Technologies, Inc.** · Berkeley, USA (Remote) *Jul 2025 – Aug 2025*

- Architected RAG-based LLM system using LangChain + FAISS serving enterprise AI agents, achieving sub-3s end-to-end query latency at production load.
- Implemented model-fallback and load-balancing logic across multi-model deployments, sustaining 99.9% uptime SLA throughout the engagement.
- Designed multi-turn conversational memory layer, reducing redundant context re-sends by ~40% and lowering token costs proportionally.
- Optimised inference pipeline for token utilisation, cutting average cost-per-query by 25% without degrading response quality.

Software Engineer · **Canara HSBC Life Insurance** · Gurugram, India *Mar 2024 – Jun 2025*

- Built Node.js microservices for authentication and onboarding flows serving millions of policy-holder sessions; added OTP throttling and retry logic that reduced failed login attempts by 35%.
- Served as Level-3 escalation owner for high-impact backend production incidents, achieving average resolution times under 2 hours on P1 issues.
- Introduced structured data profiling on onboarding datasets, identifying validation gaps that reduced form-submission errors by 20% within 60 days.
- Strengthened audit logging framework, contributing to zero compliance findings across two internal audit cycles.

Backend Developer Intern · **Paytm** · Noida, India *Sep 2021 – Nov 2021*

- Shipped 4 Node.js API endpoints and 2 React portal features adopted in the production release cycle, improving page load responsiveness by 15%.

PROJECTS

Claim Document Intelligence *Python · LangGraph · GPT-4o · FastAPI · Multi-Agent*

Apr 2026

- Architected a multi-agent LangGraph pipeline that ingests insurance claim PDFs, classifies every page into 9 document types via GPT-4o vision, then routes pages to specialist extraction agents (Identity, Discharge Summary, Itemised Bill) running in parallel.
- Each agent receives only its relevant pages (not the full PDF), minimising token usage and improving extraction accuracy; results are merged by an Aggregator node into a single structured JSON response.
- Exposed via a FastAPI service with multipart PDF upload, deterministic zero-temperature inference, and support for scanned documents — reducing manual claim data entry to a single API call.

Enterprise AI Knowledge Graph System *Python · LangChain · Neo4j · FastAPI · LLM*

Mar 2026

- Built an LLM-powered pipeline that parses unstructured technical documents, extracts entities and relationships, and stores them in Neo4j — enabling natural language graph queries through a FastAPI backend.
- Designed end-to-end architecture covering document ingestion, entity extraction, graph construction, and a Cypher-generation query engine that translates user questions into Neo4j queries via LLM.
- Exposed an interactive API and demo interface enabling retrieval of service ownership, infrastructure dependency chains, and cross-team relationships — reducing manual documentation lookups.

AI FAQ Chatbot (RAG-based Q&A) *LangChain · Pinecone · Python*

Sep – Oct 2025

- Built end-to-end document ingestion and embedding pipeline using LangChain + Pinecone with semantic search and real-time streaming, achieving <3s response latency under concurrent load.

Food Delivery Backend (RESTful APIs) *Node.js · Express.js · PostgreSQL · JWT*

Aug 2025

- Designed modular Express.js architecture with JWT-based auth and role-based access control; integrated PostgreSQL with optimised query patterns and relational schema design supporting high-frequency order and menu operations.

EDUCATION

B.Tech. in Computer Science & Engineering · Graphic Era University · 2023

CGPA: 8.87 / 10

CERTIFICATIONS

- A2A: The Agent2Agent Protocol · *Google Cloud, IBM Research · DeepLearning.AI* *Mar 2026*
- Introduction to Model Context Protocol · *Anthropic · Skilljar* *Mar 2026*
- Introduction to TensorFlow for AI, ML & Deep Learning · *DeepLearning.AI / Coursera* *Jan 2026*
- Data Visualization using Power BI *Dec 2025*
- Data Analytics & Automation: Excel, SQL, Python *Oct – Nov 2025*
- MySQL for Beginners · *Credential ID: 3097-27881460* *Dec 2022*